

Attention Is All You Need

Abstract: The dominant sequence transduction models are based on complex recurrent or convolutional neural networks that include an encoder and a decoder. The best performing models also connect the encoder and decoder through an attention mechanism. We propose a new simple network architecture, the Transformer, based solely on attention mechanisms, dispensing with recurrence and convolutions entirely.

Methodology: We used scaled dot-product attention and multi-head attention.

Key Findings: The Transformer achieves 28.4 BLEU on the WMT 2014 English-to-German translation task.